



TOXICITY PREDICTION AND EDA



Aim

I Using a dataset of 1.8million + tweets for exploring various trends in the data, and training Machine Learning and Deep Learning models to do our prediction task. Here I have treated the problem as a regression problem so that people can get the toxicity rating on a scale of 0 to 1 and not an absolute verdict on whether a message is toxic or not.

The model could also have been trained as a classification model.

Introduction

Increasing toxicity in social media messages and live streaming platforms requires automated methods that can handle these messages and report to the channel/id owners so that they can take the required action and the community can be kept clean.

Current methods mostly include manual moderation by humans which is not efficient as well as always accurate.

Currently the models are being trained on an English language dataset, and is planned to be extended to Hindi Language as well.

Exploratory Data Analysis

Finding dimensions, and primary view of the dataset.

```
train_df.shape
```

```
(1804874, 45)
```

Thus, we can see that the dataset has 45 columns (indicating 45 features) and 1.8 million+ tuples , indicating 1.8m+ tweets.

This is the primary view of the dataset:

```
#describe the dataset  
train_df.head(5)
```

	id	target	comment_text	severe_toxicity	obscene	identity_attack	insult	threat	asian	atheist	...	article_id	rating	fu
0	59848	0.000000	This is so cool. It's like, 'would you want yo...	0.000000	0.0	0.000000	0.00000	0.0	NaN	NaN	...	2006	rejected	
1	59849	0.000000	Thank you!! This would make my life a lot less...	0.000000	0.0	0.000000	0.00000	0.0	NaN	NaN	...	2006	rejected	
2	59852	0.000000	This is such an urgent design problem; kudos t...	0.000000	0.0	0.000000	0.00000	0.0	NaN	NaN	...	2006	rejected	
3	59855	0.000000	Is this something I'll be able to install on m...	0.000000	0.0	0.000000	0.00000	0.0	NaN	NaN	...	2006	rejected	
4	59856	0.893617	haha you guys are a bunch of losers.	0.021277	0.0	0.021277	0.87234	0.0	0.0	0.0	...	2006	rejected	

5 rows × 45 columns

Some of the columns in the dataset are: (Most of them are numbers).

0	id	int64
1	target	float64
2	comment_text	object
3	severe_toxicity	float64
4	obscene	float64
5	identity_attack	float64
6	insult	float64
7	threat	float64
8	asian	float64
9	atheist	float64
10	bisexual	float64
11	black	float64
12	buddhist	float64
13	christian	float64
14	female	float64
15	heterosexual	float64
16	hindu	float64
17	homosexual_gay_or_lesbian	float64
18	intellectual_or_learning_disability	float64
19	jewish	float64
20	latino	float64
21	male	float64
22	muslim	float64
23	other_disability	float64
24	other_gender	float64
25	other_race_or_ethnicity	float64
26	other_religion	float64
27	other_sexual_orientation	float64
28	physical_disability	float64
29	psychiatric_or_mental_illness	float64
30	transgender	float64
31	white	float64
32	created_date	object
33	publication_id	int64
34	parent_id	float64
35	article_id	int64
36	rating	object
37	funny	int64
38	wow	int64

Number of Parent tweets and number of tweet replies distribution in the dataset: (1 is for parent, 0 is for reply).

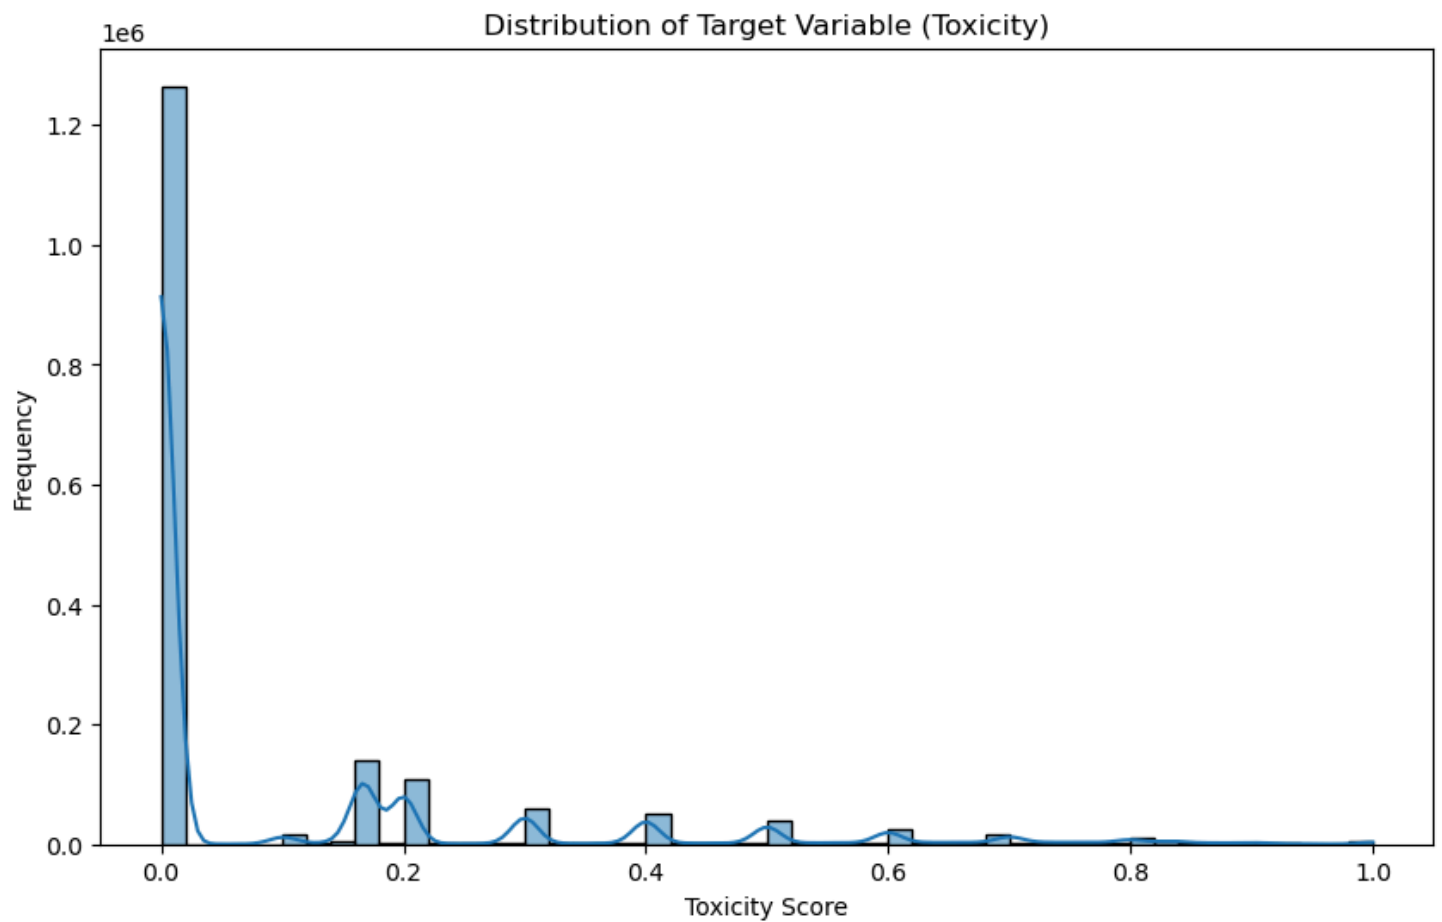
```
is_reply
1      1026225
0       778646
Name: count, dtype: int64
```

Analysing the target column:

Number of unique values of the target column:

```
target
0.000000    1264761
0.166667     138501
0.200000     107492
0.300000      59098
0.400000      50013
...
0.026684         1
0.924561         1
0.007458         1
0.145161         1
0.870088         1
```

Distribution of toxicity rating across the dataset:




Thus, we can see that most of the comments are non-toxic (which have a toxicity score of 0)

Percentage of Toxic and Non-Toxic comments when considered as a classification problem:

(Considering the threshold for being toxic as 0.5)

Count of non-toxic and toxic comments from a classification PoV



Thus we can see that the dataset we are using is not uniform and mostly skewed with most of the comments being non toxic as we saw earlier.

Looking for sub-features:


There are several sub-classifications of a comment if it is found out to be toxic.

1. severe_toxicity
2. obscene
3. identity_attack
4. insult
5. threat

Their datatypes:



1	target	float64
2	comment_text	object
3	severe_toxicity	float64
4	obscene	float64
5	identity_attack	float64
6	insult	float64
7	threat	float64

Finding correlation among these sub-features and the target variable:



Thus, we can see that all of them are positively correlated, since all of them have values > 0 .

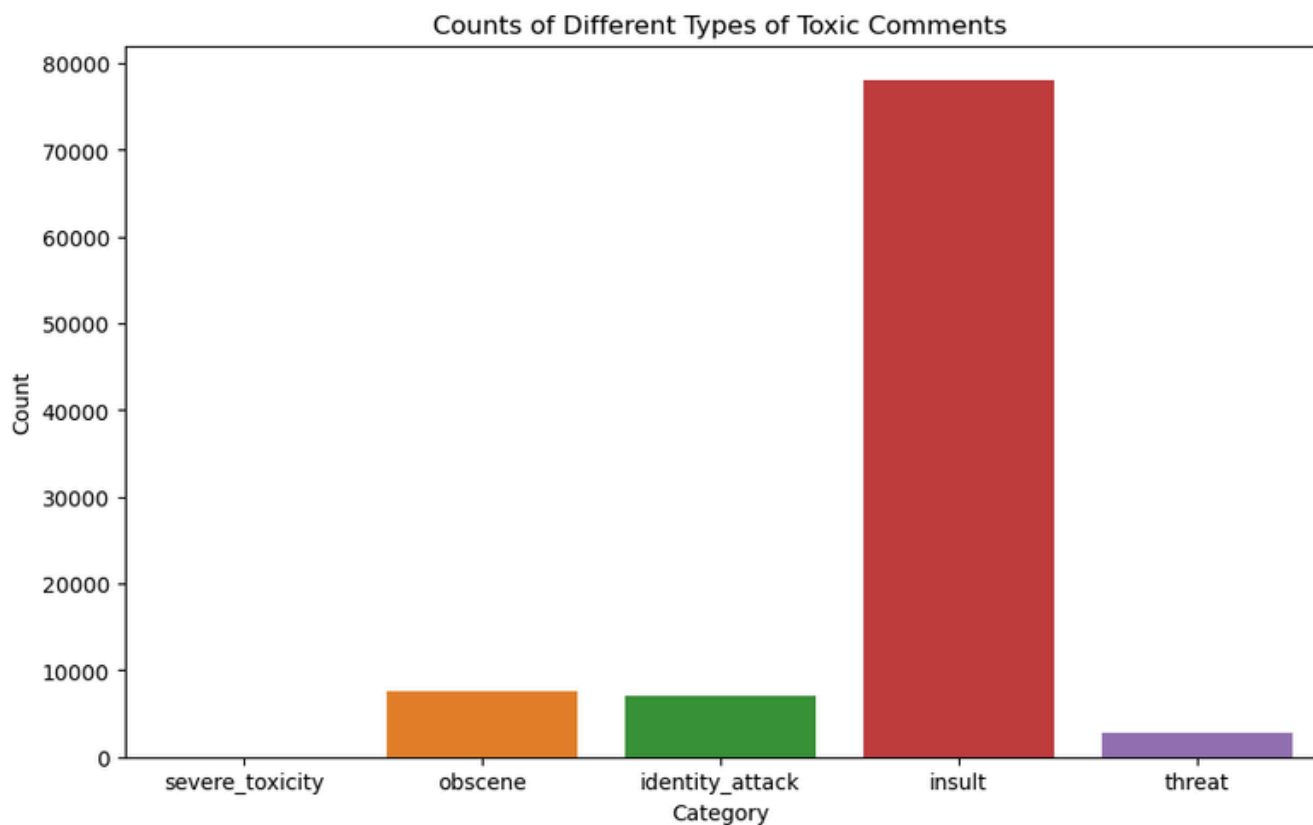
Their distribution in the Entire Dataset, and Toxic comments Dataset:



Thus we can see, that most of the comments have their ratings close to zero, and very few have higher ratings.

Among the toxic comments, most of them are insults.

Cross verification by plotting bar plots (considering 0.5 to be the threshold for each of these subclasses):




Our observation that most of the toxic comments are insults turned out to be correct.

Several other Features are present in the dataset:

1. Gender related
2. Sexual orientation related
3. Race/Ethnicity related


Let us try to analyse which kind of reference hints towards which kind of toxicity:

Percentage of type of toxicity in comments where gender references are made




Thus we can see that toxic comments that include gender references are mostly Insults in nature.

Percentage of type of toxicity in comments where sexual orientation references are made



Thus we can see that mostly the toxic comments where sexual orientation references are made come under the category of identity_attack.


Percentage of type of toxicity in comments where race/ethnic references are made



We can observe that toxic comments that include racial/ethnicity comments mostly fall under the category of identity attack or insult.

Observing the distribution of these values show that most of the comments do not have references of these kind and hence point towards zero, so we do not plot this distribution for other similar features as we cannot get any major insights from them.

Distribution of gender values in the comments



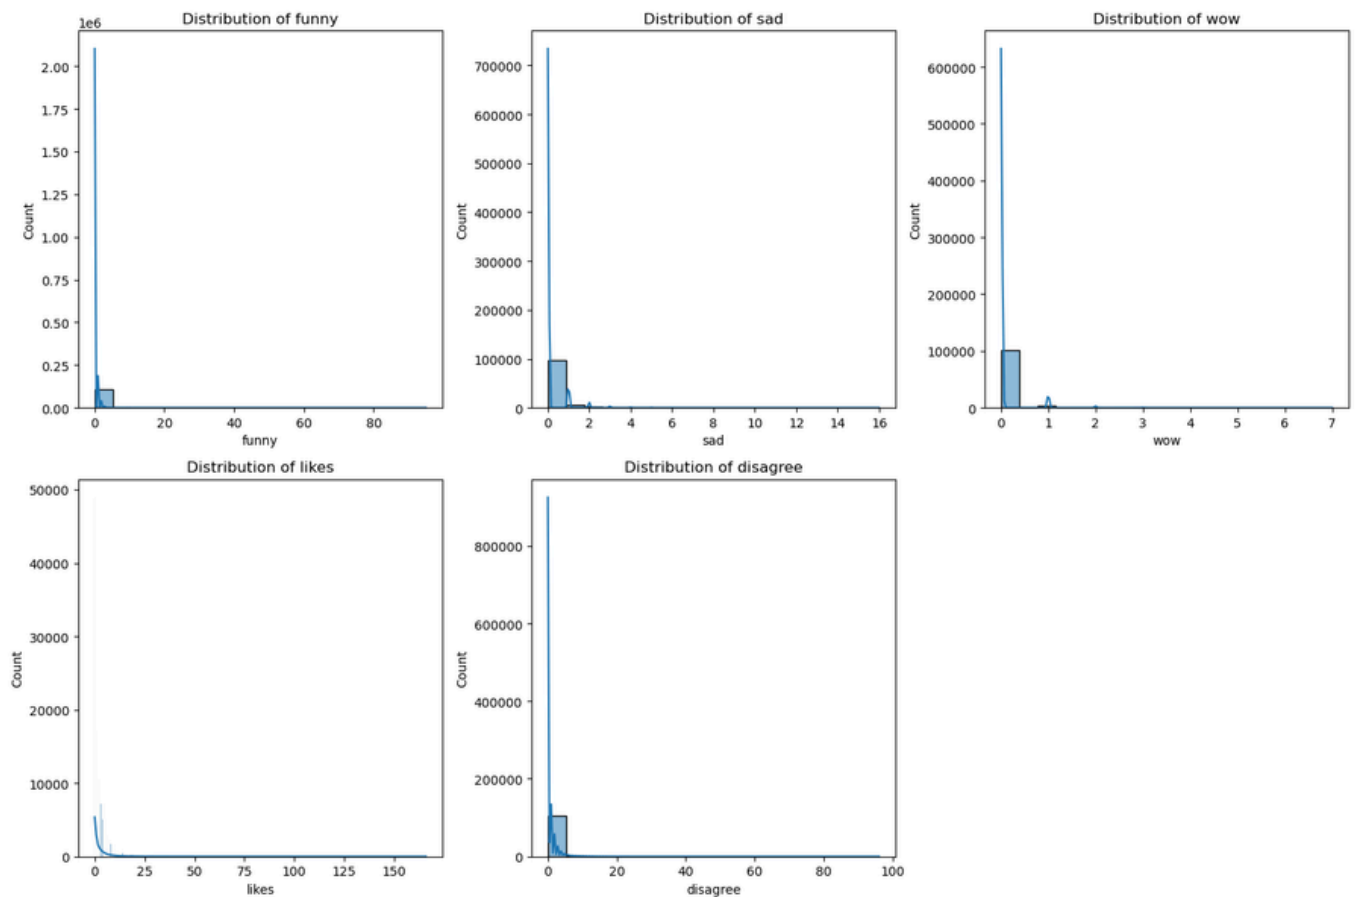
User feedback Features:

- funny
- sad
- wow
- likes
- disagree

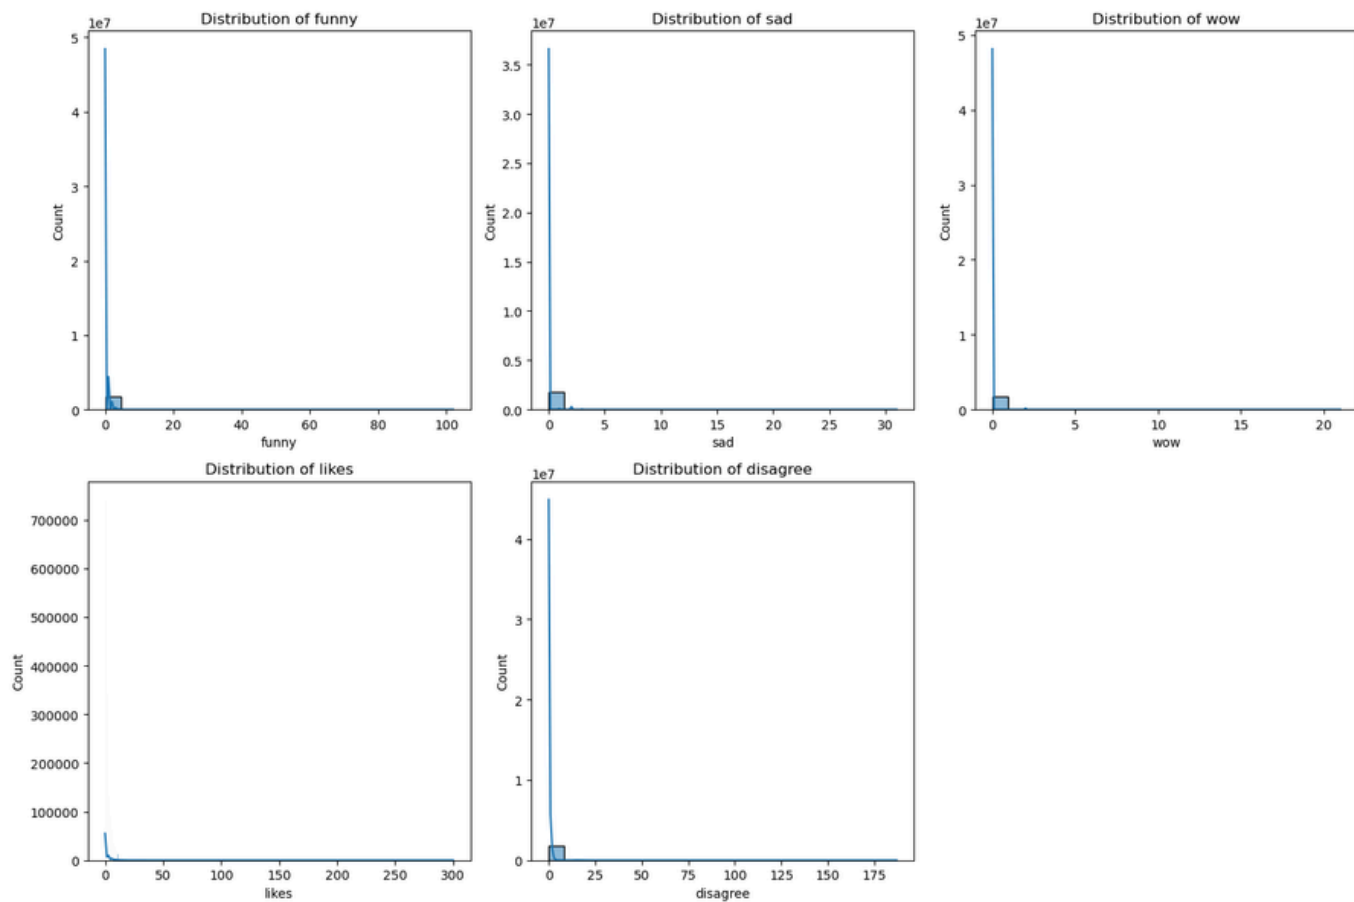
Their statistics:

	funny	sad	wow	likes	disagree
count	1.804871e+06	1.804871e+06	1.804871e+06	1.804871e+06	1.804871e+06
mean	2.779246e-01	1.091175e-01	4.420704e-02	2.446166e+00	5.843692e-01
std	1.055308e+00	4.555366e-01	2.449361e-01	4.727925e+00	1.866590e+00
min	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
25%	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00	0.000000e+00
50%	0.000000e+00	0.000000e+00	0.000000e+00	1.000000e+00	0.000000e+00
75%	0.000000e+00	0.000000e+00	0.000000e+00	3.000000e+00	0.000000e+00
max	1.020000e+02	3.100000e+01	2.100000e+01	3.000000e+02	1.870000e+02

Their distribution in toxic comments:



Their distribution in overall dataset:



Since most of the values are close to 0 again(for toxic as well as non toxic comments) , there is no valuable insights that can be gained from these features.

Some more statistical Explorations:

Average length of toxic comments: 252.8707510475582

Average length of non-toxic comments: 300.01503680156947

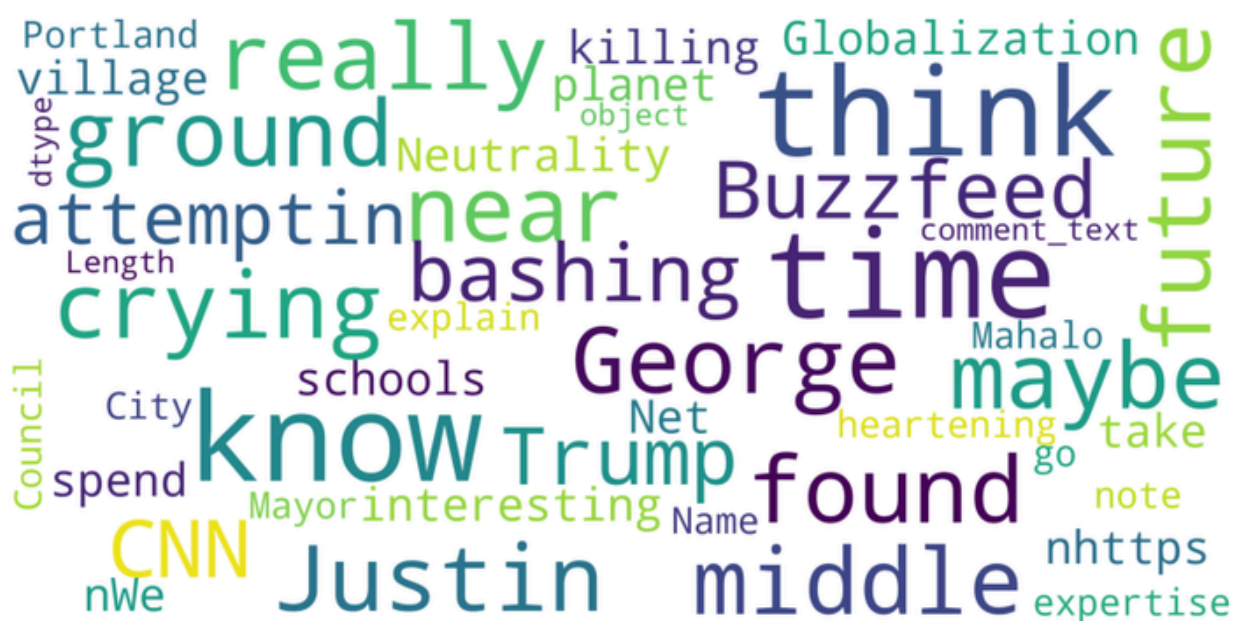
Average number of exclamations in toxic comments :0.3984948984385276

Average number of exclamations in non-toxic comments : 0.21415799151335377

Toxic comments are being noticed to be slightly shorter than non toxic comments.

Also, toxic comments have more use of special character ‘!’ and most probably other special characters too.

Vizualising prevalent words using WordCloud:



Prevalent words in all comments



Prevalent words in toxic comments





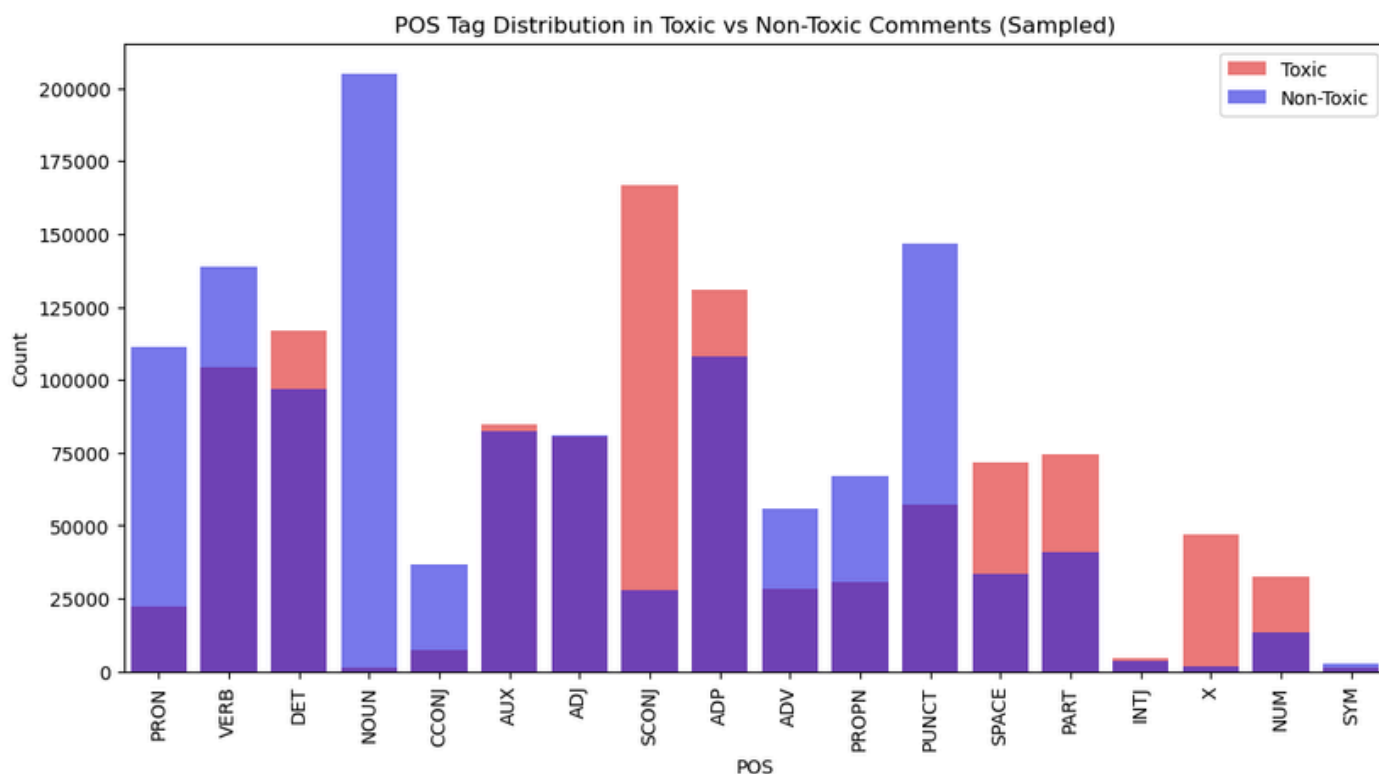
Prevalent words in severe toxic comments



Prevalent words in clean/non-toxic comments

Note: These wordclouds have been created from a sample of the data due to computational resource limitations. They are good for providing a generalised idea but may not be completely accurate.

PoS Tagging for additional Insights:



- Although the data used here for POS Tagging is random sampled, I'm pretty sure that the distribution would change if the entire set of toxic and non toxic comments was considered for this analysis.
- However, one major analysis which I think is valid is that toxic comments are more likely to contain words that do not fall into any specific category (denoted by X), probably because toxic messages are more likely to contain short forms, slang words, typos and so on which may not always be grammatically correct.

Regression Algorithms

Data Preprocessing:

Handling missing values- Data without actual text is of no use to us.

```
#removing rows that have the text missing
train_df.dropna(subset=['comment_text'], inplace=True)
```

Lemmatization and removal of Stopwords:

((I have used Spacy library to achieve this)

```
def preprocess(text_string):
    text_string = text_string.lower()
    text_string = re.sub(r'^A-Za-z0-9+', ' ', text_string) #Remove special characters and punctuations
    doc = nlp(text_string)
    new_text = []
    for token in doc:
        if token.text.lower() not in STOP_WORDS:
            new_text.append(token.lemma_)
    text_string = ' '.join(new_text)
    return text_string
```

Training the models:

1. Bag of Words (BoW) approach:

a. SGD Regressor (Stochastic Gradient Descent Regressor)

```
#Hyperparameters tuning
param_grid={ 'alpha' :[0.00001, 0.0001, 0.001, 0.01, 0.1, 1, 10, 100],
              'penalty' :['l1', 'l2']}
sgd=SGDRegressor()
```

Best Parameters: {'alpha': 0.0001, 'penalty': 'l1'}
Mean Squared Error on CV set: 0.024298003603473017

b. Decision Tree Regressor

```
param_grid = {'max_depth': [3, 5, 7],
               'min_samples_leaf': [10, 100, 1000]}

dt=DecisionTreeRegressor()
```

Best Parameters: {'max_depth': 7, 'min_samples_leaf': 100}
Mean Squared Error on CV set: 0.031656608167548304

Feature weights in SGDreg

	weights
idiot	0.434650
stupid	0.350389
stupidity	0.323367
moron	0.320754
pathetic	0.301164
hypocrite	0.283425
crap	0.266751
dumb	0.266327
idiotic	0.264953
ass	0.246647
ignorant	0.234407
clown	0.228581
damn	0.215632
scum	0.211157
ridiculous	0.208342
fool	0.206690
loser	0.205507
jerk	0.198379
silly	0.197860
shit	0.184310

Feature weights in decision Tree

	weights
stupid	0.354878
idiot	0.257678
ignorant	0.086890
pathetic	0.074313
dumb	0.072259
fool	0.069534
stupidity	0.065688
year	0.002779
don fool	0.002214
people	0.002021
don	0.001853
good	0.001797
thing	0.001514
state	0.001183
time	0.000840
way	0.000754
go	0.000627
need	0.000608
issue	0.000584
public	0.000438

2. TF-IDF approach:

a. SGD Regressor (Stochastic Gradient Descent Regressor)

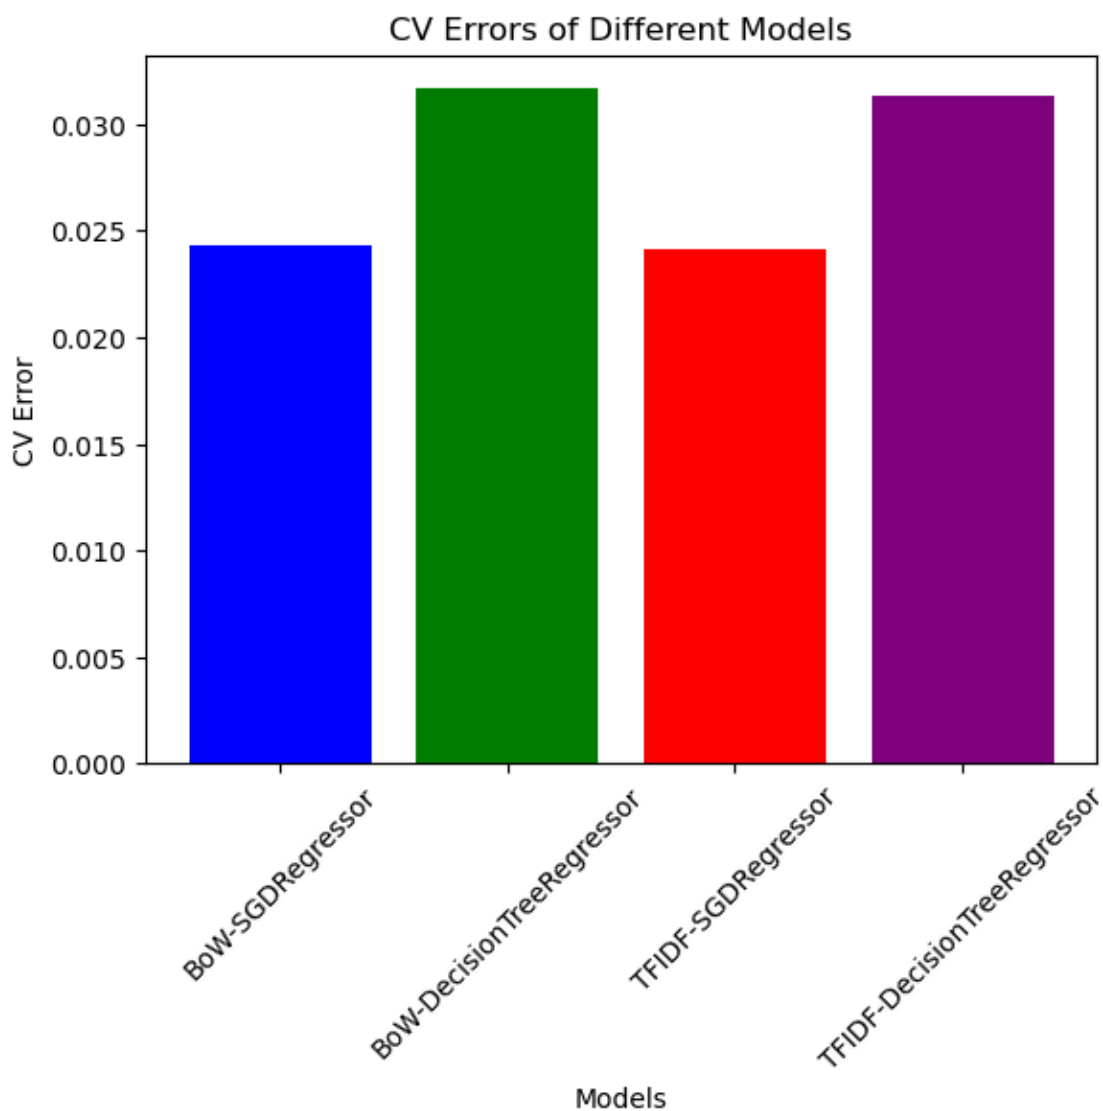
Best Parameters: {'alpha': 1e-05, 'penalty': 'l2'}

Mean Squared Error on CV set: 0.02415963883069067

b. Decision Tree Regressor


Best Parameters: {'max_depth': 7, 'min_samples_leaf': 10}

Mean Squared Error on CV set: 0.03138149202136317



LSTM

I used two LSTM layers and a dense layer, output function used was Sigmoid since we want values from 0 to 1, and optimizer used was rmsprop.



Outcomes of LSTM Model:

- LSTM Model : Mean Squared Error on CV set: 0.0157

This was the best model trained so far with minimum Mean Squared Error on cross validation set.